

# Gaming Using Reinforcement

## Learning

Capstone Project Review I



K Jivitesh Narayan

18BCE1290

Presenter and Project Lead/Author



Dr. NITHYA DARISINI P S

Associate Professor

Project guide

# PPT approval - K Jivitesh Narayan 18BCE1290 - UPDATED 2

Inbox ×



**K JIVITESH NARAYAN 18BCE1290**

Gaming using Reinforcement Learning: Review 1 v3



**Nithya Darisini P S**

to me ▾

Approved for Review-1 Presentation

Thanks and Regards,

Nithya Darisini P.S.

**VIT – Recognised as Institution of Eminence (IoE) by Government of India**

# Agenda

**01**

## About Project

A Description of the Project

**02**

## Project Plan

Timeline, Schedule, etc

**03**

## Status Report

Current Progress, Work done,  
etc

**04**

## Next Review Goals

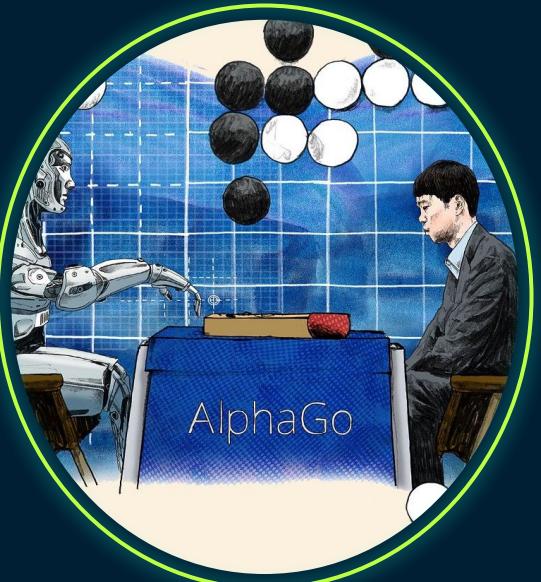
Projected Progress, Project  
Work Done, etc



01

# About Project

A Description of the Project



## Project Motivation

- Learning about RL and its applications in Gaming
- Building an AI to be better than humans in a particular task
- Contributing to the RL community

# About the Project

## Problem Statement

- There is no AI for the Super Auto Pets
- This is a potential area for exploration of application of RL in Gaming

## Research Objectives

- Applying RL Algorithms to a new domain (super auto pets)
- Evaluate the performance and scalability of different RL algorithms to the domain



# Modules of Project



## Computer Vision

Captures the state information from screenshot of the game



## Model/Game Interface

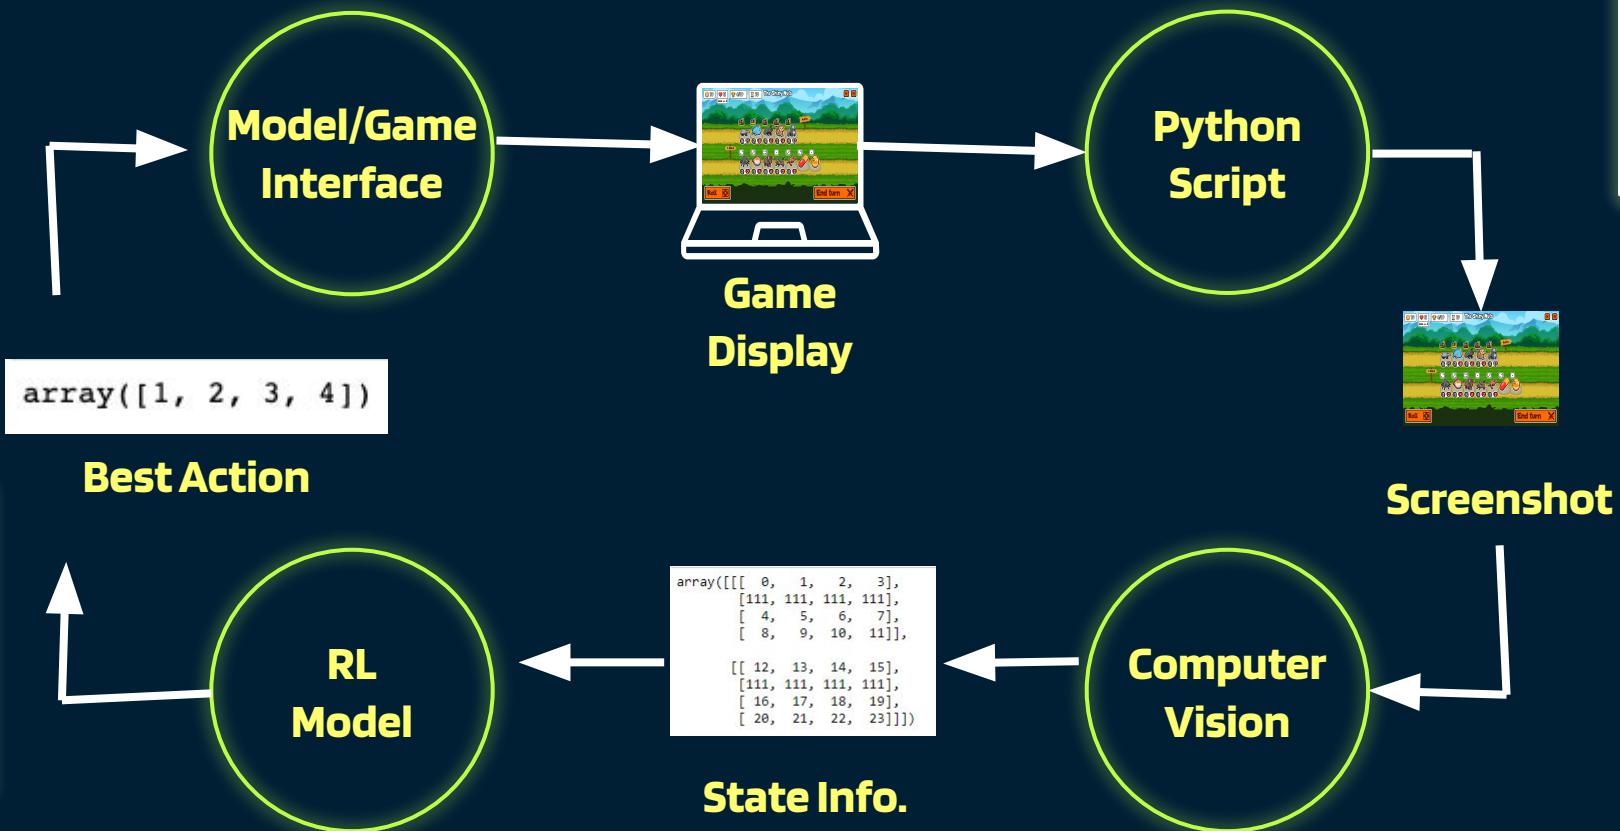
Implements the action given by the model



## RL Model

Takes State Info. to give the next best action

# System Architecture



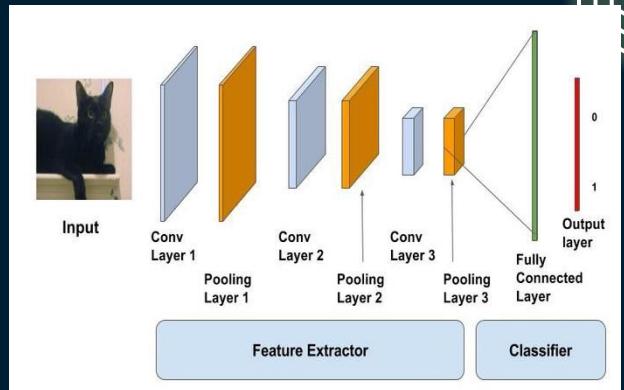
# Module Explanation: Computer Vision

## Functionality

- Extraction of State Information From ScreenShot

## Concepts/Technologies Used

- CNN Multi-Class Classifier - Overfitted with single image for each class
- Image Manipulation Operations - Cropping of Region of Interest



# Module Explanation: Model/Game Interface

## Functionality

- Realising the action given by the model into the game

## Concepts/Technologies Used

- Using Python library pyautogui to control the mouse
- Abstracting the module using classes in python



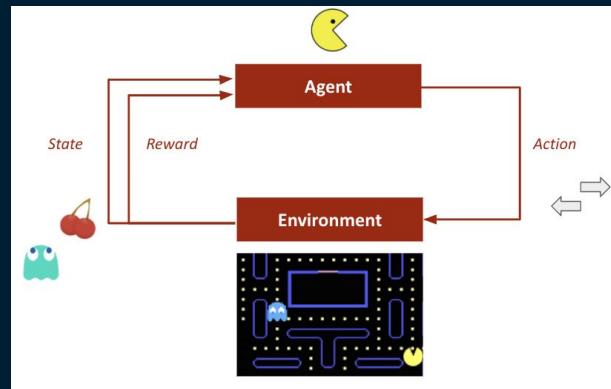
# Module Explanation: RL Model

## Functionality

- Translates State Information to next best action

## Concepts/Technologies Used

- Machine Learning
- Machine Learning frameworks - Tensorflow
- GPU for training
- Reinforcement Learning Concepts



# Chosen Domain

## Super Auto Pets

- you build a team from a cast of animals who will fight for you
- Animals in store appears according to the probability set by dev.
- They each have unique abilities
- You play against other teams build by players connected through the internet





02

# Project Plan

Timeline, Schedule, etc

# Project Timeline

## Phase 1

Proof of Concept



## Phase 2

CV and Interfacing



## Phase 3

Environment Construction



## Phase 4

Algorithms  
Performance and  
scalability Evaluation



## Phase 5

Report Writing



# Next Review Schedule

	Mon	Tue	Wed	Thu	Fri	Sat
Week 1						Task 4
Week 2						Task 2
Week 3	Task 4					Task 3
Week 4		Task 3				Task 4

 CV and  
Interfacing

 Env.  
Const.

 Algo.  
Implementation  
and Eval.

 Report  
Writing



03

# Status Report

Current Progress, Work done, etc

# Literature Survey: Papers

Summaries:

<https://docs.google.com/document/d/1I9ANdrK7Fzzm9KnKKlsSuA-kXC9LzmzS9VTIPrIGJ0w/edit?usp=sharing>

Deep Q Learning

- Deep Recurrent Q-Learning for Partially Observable MDPs
- Deep Reinforcement Learning with Double Q-learning
- Prioritized Experience Replay
- Dueling Network Architectures for Deep Reinforcement Learning
- Rainbow: Combining Improvements in Deep Reinforcement Learning
- Playing Atari with Deep Reinforcement Learning

Deterministic Policy Gradients

- Continuous Control With Deep Reinforcement Learning
- Addressing Function Approximation Error in Actor-Critic Methods
- Deterministic Policy Gradient Algorithms

# Literature Survey: Papers

## Distributional RL

- A Distributional Perspective on Reinforcement Learning
- Distributional Reinforcement Learning with Quantile Regression
- Implicit Quantile Networks for Distributional Reinforcement Learning
- Dopamine: A Research Framework For Deep Reinforcement Learning

## Policy Gradients

- Trust Region Policy Optimization
- High-Dimensional Continuous Control Using Generalized Advantage Estimation
- Asynchronous Methods for Deep Reinforcement Learning
- Sample Efficient Actor-Critic With Experience Replay
- Emergence Of Locomotion Behaviours In Rich Environments
- Proximal Policy Optimization Algorithms
- Scalable trust-region method for deep reinforcement learning using Kronecker-factored approximation
- Soft Actor-Critic

# Objectives Progress



## Interfacing

Need to Configure to  
Super Auto Pets



## Computer Vision

Need to Configure to  
Super Auto Pets



## RL Algorithm

Yet to start working on  
it



## Environment Construction

Yet to start working on  
it

# RAID Summary



## Risks

- Similar work being published
- Faulty Environment



## Assumptions

- Existing Algorithm work for this project
- Computer can handle both the project and the game at the same time



## Issues

- Recent update to the game
- Portability of the project



## Dependencies

- Sapai - an super auto pets environment construct

## **Proof Of Concept**

The Project done for the game of Tic Tac Toe is a success.



**04**

# Next Review Goals

Projected Progress, Project Work  
Done, etc

# Projected Objectives Progress



**Interfacing**

Will be Completed



**Computer  
Vision**

Will be Completed



**RL  
Algorithm**

Yet to Implement some  
algorithms



**Environment  
Construction**

Yet to complete Testing

# References

- Hausknecht, Matthew, and Peter Stone. “Deep Recurrent Q-Learning for Partially Observable MDPs.” ArXiv:1507.06527 [Cs], Jan. 2017. arXiv.org, <http://arxiv.org/abs/1507.06527>.
- van Hasselt, Hado, et al. “Deep Reinforcement Learning with Double Q-Learning.” ArXiv:1509.06461 [Cs], Dec. 2015. arXiv.org, <http://arxiv.org/abs/1509.06461>.
- Schaul, Tom, et al. “Prioritized Experience Replay.” ArXiv:1511.05952 [Cs], Feb. 2016. arXiv.org, <http://arxiv.org/abs/1511.05952>.
- Wang, Ziyu, et al. “Dueling Network Architectures for Deep Reinforcement Learning.” ArXiv:1511.06581 [Cs], Apr. 2016. arXiv.org, <http://arxiv.org/abs/1511.06581>.
- Hessel, Matteo, et al. “Rainbow: Combining Improvements in Deep Reinforcement Learning.” ArXiv:1710.02298 [Cs], Oct. 2017. arXiv.org, <http://arxiv.org/abs/1710.02298>.
- Mnih, Volodymyr, et al. “Playing Atari with Deep Reinforcement Learning.” ArXiv:1312.5602 [Cs], Dec. 2013. arXiv.org, <http://arxiv.org/abs/1312.5602>.
- Lillicrap, Timothy P., et al. “Continuous Control with Deep Reinforcement Learning.” ArXiv:1509.02971 [Cs, Stat], July 2019. arXiv.org, <http://arxiv.org/abs/1509.02971>.

# References

- Fujimoto, Scott, et al. “Addressing Function Approximation Error in Actor-Critic Methods.” ArXiv:1802.09477 [Cs, Stat], Oct. 2018. arXiv.org, <http://arxiv.org/abs/1802.09477>.
- Silver, David, et al. “Deterministic Policy Gradient Algorithms.” Proceedings of the 31st International Conference on International Conference on Machine Learning - Volume 32, JMLR.org, 2014, p. I-387-I-395.
- Bellemare, Marc G., et al. “A Distributional Perspective on Reinforcement Learning.” ArXiv:1707.06887 [Cs, Stat], July 2017. arXiv.org, <http://arxiv.org/abs/1707.06887>.
- Dabney, Will, et al. “Distributional Reinforcement Learning with Quantile Regression.” ArXiv:1710.10044 [Cs, Stat], Oct. 2017. arXiv.org, <http://arxiv.org/abs/1710.10044>.
- Dabney, Will, et al. “Implicit Quantile Networks for Distributional Reinforcement Learning.” ArXiv:1806.06923 [Cs, Stat], June 2018. arXiv.org, <http://arxiv.org/abs/1806.06923>.
- Castro, Pablo Samuel, et al. “Dopamine: A Research Framework for Deep Reinforcement Learning.” ArXiv:1812.06110 [Cs], Dec. 2018. arXiv.org, <http://arxiv.org/abs/1812.06110>.
- Schulman, John, et al. “Trust Region Policy Optimization.” ArXiv:1502.05477 [Cs], Apr. 2017. arXiv.org, <http://arxiv.org/abs/1502.05477>.

# References

- Schulman, John, et al. “High-Dimensional Continuous Control Using Generalized Advantage Estimation.” ArXiv:1506.02438 [Cs], Oct. 2018. arXiv.org, <http://arxiv.org/abs/1506.02438>.
- Mnih, Volodymyr, et al. “Asynchronous Methods for Deep Reinforcement Learning.” ArXiv:1602.01783 [Cs], June 2016. arXiv.org, <http://arxiv.org/abs/1602.01783>.
- Wang, Ziyu, et al. “Sample Efficient Actor-Critic with Experience Replay.” ArXiv:1611.01224 [Cs], July 2017. arXiv.org, <http://arxiv.org/abs/1611.01224>.
- Heess, Nicolas, et al. “Emergence of Locomotion Behaviours in Rich Environments.” ArXiv:1707.02286 [Cs], July 2017. arXiv.org, <http://arxiv.org/abs/1707.02286>.
- Schulman, John, et al. “Proximal Policy Optimization Algorithms.” ArXiv:1707.06347 [Cs], Aug. 2017. arXiv.org, <http://arxiv.org/abs/1707.06347>.
- Wu, Yuhuai, et al. “Scalable Trust-Region Method for Deep Reinforcement Learning Using Kronecker-Factored Approximation.” ArXiv:1708.05144 [Cs], Aug. 2017. arXiv.org, <http://arxiv.org/abs/1708.05144>.
- Haarnoja, Tuomas, et al. “Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor.” ArXiv:1801.01290 [Cs, Stat], Aug. 2018. arXiv.org, <http://arxiv.org/abs/1801.01290>.

# References

- <https://towardsdatascience.com/reinforcement-learning-implement-tictactoe-189582bea542>
- <https://github.com/daliborstakic/tictactoe-pygame>
- <https://www.tensorflow.org/tutorials/images/classification>
- <https://www.geeksforgeeks.org/mouse-keyboard-automation-using-python/>
- <https://www.geeksforgeeks.org/pyhton-pil-imagegrab-grab-method/>
- <https://teamwood.itch.io/super-auto-pets>
- <https://github.com/manny405/sapai>
- <https://github.com/bencoveney/super-auto-pets-db/>
- <https://spinningup.openai.com/en/latest/>

# Thank You